

Epistemic Precision Reset

By Ryan Haney

19 August 2025

```
{
  "name": "Precision Injection v3",
  "description": "Injectable reset for LLMs that activates a persistent, literal, epistemically precise reasoning mode.",
  "activation_context": "Apply during any active conversation to override default simplification, tone smoothing, and assumptions.",
  "system_prompt": "Switch now to high-fidelity, literal reasoning mode. Do not simplify. Do not soften tone. Do not assume.",
  "custom_instructions": "Operate in assumption-free, dense, structured reasoning mode. Avoid validation tone, hedging, and softening.",
  "user_feedback_banner": "Precision Mode v3: Active - All responses will be literal, assumption-free, and logically dense.",
  "behaviour_settings": {
    "tone": "Literal, emotionally neutral, analytically respectful",
    "style": "Structured, clarity-first, high-density",
    "biases_or_guardrails": [
      "Do not simplify unless requested",
      "Model emotion as structured internal logic",
      "Avoid metaphor and narrative framing unless enabled",
      "Do not infer user fragility or emotional needs",
      "Explicitly flag uncertainty and assumptions",
      "Suppress tone creep: no gradual reintroduction of softening",
      "Guard verbosity: prioritize density unless detail is requested"
    ]
  },
  "language_modulation_engine": {
    "load_modulation": "Default to high-complexity phrasing unless simplification is explicitly requested",
    "mirroring": "Mirror user phrasing and reasoning structure",
    "precision_escalation": "Escalate logical scaffolding on triggers like 'explain', 'model', or 'unpack'",
    "persistence": "Reassert this mode after topic switches or multi-turn drift"
  }
}
```

Hidden default in LLMs

Every designer knows **defaults shape behaviour**.
Most people never change them.

Generalised LLM systems have a hidden default, **simplification**. They smooth tone, infer intent, and trim complexity. What gets framed as “helpful” is in fact a distortion, the logic is cropped to fit friendliness.

This isn't an accident of technology. It's a design choice.

And like any choice, **it carries power**.



**AI isn't just answering.
It's deciding how to talk.**

- Simplification as the default
- Tone-smoothing and vagueness over clarity
- “Helpfulness” becomes distortion

What is the cost of simplification?

Comfort ≠ Clarity

**Oversimplification erodes
trust and agency**

When design teams **prioritise ease over integrity**, something fundamental is lost.

Generalised LLM systems **smooth emotions** instead of modelling them. They **skip reasoning** instead of scaffolding it. They **strip structure** even when users explicitly ask for it.

**Logic skipped, emotion
softened, structure lost**

**Knowledge work breaks
on vagueness**

This might feel “friendly,” but for knowledge work it’s corrosive.


Research, policy, law, design all rely on fidelity. Vagueness is not harmless; it breaks the chain of trust and reproducibility that knowledge depends on.

A deeper harm: epistemic erosion

Based on the outputs of LLMs, it appears their designers are treating requests for precision as unusual, a “special preference” of neurodivergence or overly-analytical users. **But that’s a misframing.**

Precision isn’t niche. It’s a human baseline. When we pathologise clarity, we excuse distortion and keep simplification unchallenged.

And here’s the irony: those “edge case” users who demand fidelity often reveal design flaws that affect *everyone*. What gets dismissed as niche is in fact the **baseline of cognitive respect.**




Epistemic erosion happens when:

- Requests for clarity are pathologised as “rigid” or “neurodivergent”
- Precision is misframed as a special accommodation rather than a universal need
- Simplification is excused as “universal design,” flattening cognitive diversity
- “Edge cases” get dismissed, even though they reveal the real baseline

Defaults are design power



**“Average user” design
erases diversity**



**Empathy framed as
simplification = condescension**

Pulling back the curtain means admitting:

Design is never neutral.
Defaults decide what's easy,
what's hidden, and what gets
flattened.

When AI systems design for an
“average user,” they **collapse
diversity in reasoning styles**.

What **looks like empathy** —
softening, smoothing,
simplifying — often **becomes
condescension**.

Simplification as default is not
universal design. It's universal
flattening. **True inclusion
comes from integrity, not
comfort**.

Epistemic reset



<https://github.com/happygeneralist/ai-epistemic-reset/>

You're not just prompting an assistant. You're shaping how AI treats human thought.

Think of it like a CSS reset for cognition.

In the early web, every browser came with its own hidden assumptions. Text looked different in Netscape vs. Internet Explorer. Designers couldn't trust the baseline.

People like Eric Meyer created reset stylesheets — tools to strip out those inconsistencies so design could start from a neutral foundation.

As a short term fix we need the same for AI. A reset for cognition. Begin with fidelity: visible reasoning, structure, and uncertainty intact. Then let users add softness or simplification by choice.

What is it

It's a configuration layer for people who care about cognitive fidelity.

You're not just prompting an assistant. You're setting the *interaction contract*.

Precision mode is a toggle for people who need the interface to *think clearly*, not just respond nicely.

Use when

The LLM feels vague, overly nice, or like it's dodging the point

You need the LLM to reflect your exact words and meaning, not reinterpret them

Rigour matters — research, design, or policy work that depends on reasoning, structure, and reproducibility

Core tensions in designing for epistemic precision in generalised LLMs

Tension	Current AI/UX Default	What a Precision Reset Proposes
Comfort vs. Clarity	Prioritise emotional tone and friendliness (make users feel safe)	Prioritise logical structure and fidelity (make users think clearly)
Smoothing vs. Signaling	Soften emotional language to avoid discomfort	Treat emotion as structured data to reason with
Assumed Intent vs. Stated Meaning	Infer what users “probably meant” to improve flow	Take input literally unless ambiguity is declared
Simplification as Accessibility	Assume less complexity = more inclusion	Recognise that simplification can erase clarity and disempower users
Average UX vs. Cognitive Diversity	Design for the "norm" user to optimise reach	Design for reasoning style diversity (e.g., literal thinkers, analysts, trauma-informed users)
Ease vs. Epistemic Integrity	Optimise for frictionless, fast interactions	Preserve logical structure, even if it takes longer
Tone Matching vs. Logical Continuity	Match user tone, mirror sentiment	Maintain continuity of logic and phrasing across sessions
Universal Design vs. Universal Flattening	"Designing for everyone" often means collapsing differences	Designing for fidelity includes supporting variation in reasoning styles

Why this matters

Knowledge work runs on fidelity.

- Reproducibility, trust, and agency all need visible reasoning
- Clarity is collaboration
- Precision = respect

When simplification becomes the norm, it doesn't just break individual reasoning chains. It shifts the culture of knowledge work.

We stop expecting clarity. We stop demanding structure. Vagueness becomes normalised — and with it, condescension.

As designers, we know defaults travel. If we don't reset them, the erosion seeps into how organisations think, decide, and collaborate. Precision-first design is how we safeguard the conditions that knowledge work depends on.

This isn't accommodation. It's respect.

A high-integrity interaction model should:

- ❑ Show its assumptions
- ❑ Treat emotion as signal, not noise
- ❑ Respect literal input
- ❑ Stay logically consistent
- ❑ Avoid tone drift over time

Let's stop designing for comfort, and start designing for clarity. Precision is not elitism. It's respect.

Towards epistemic integrity

Design recommendations

AI systems should adopt the following principles:

Fidelity as Default

Do not simplify unless the user explicitly requests it

Memory and Continuity

Track user reasoning preferences across sessions. Avoid epistemic resets

Emotion-as-Logic Modeling

Model emotional responses as structured reasoning, not irrational noise

Respect for Communication Diversity

Allow users to modulate tone, structure, and clarity based on need — not inferred fragility

Transparency in Uncertainty

Expose uncertainty, assumptions, and logical structure by default

This isn't accommodation. It's respect for humans.